Exercise 2

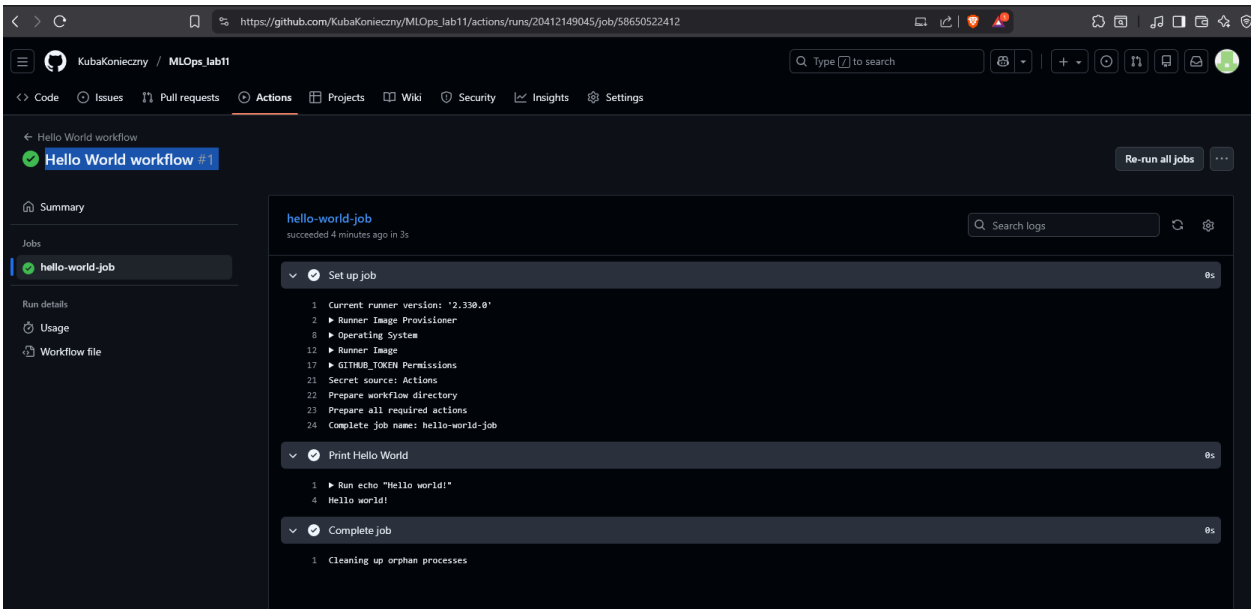

Exercise 3,4 and Homework

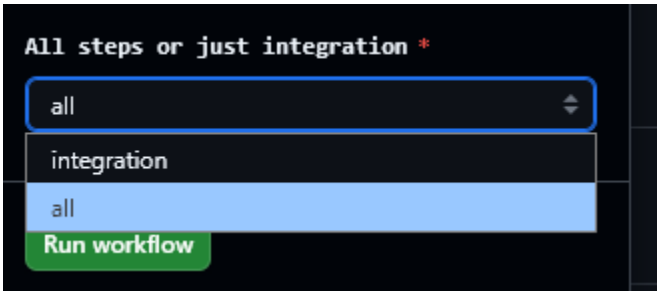

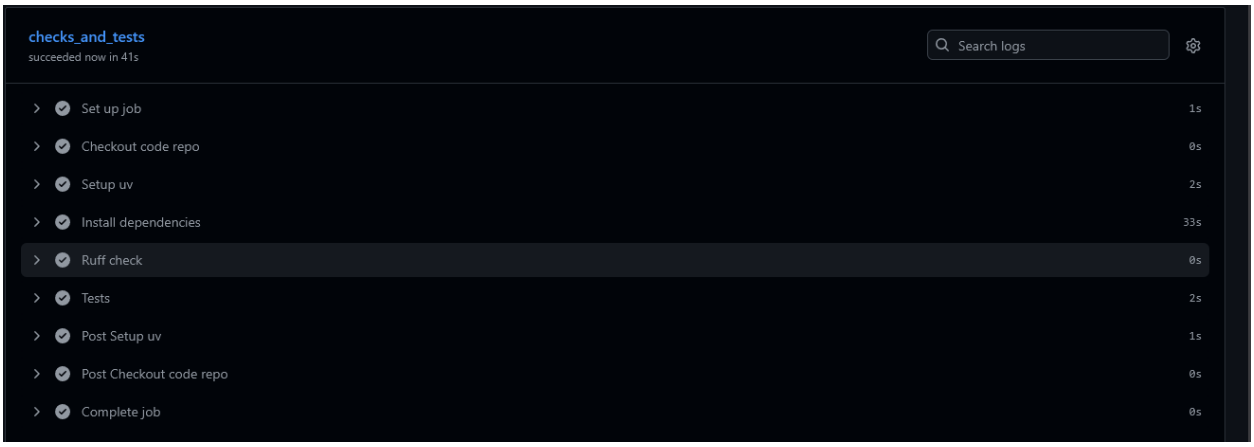
checks_and_tests
succeeded now in 41s
Search logs
Set up job 1s
Checkout code repo 0s
Setup uv 2s
Install dependencies 33s
Ruff check 0s
Tests 2s
Post Setup uv 1s
Post Checkout code repo 0s
Complete job 0s

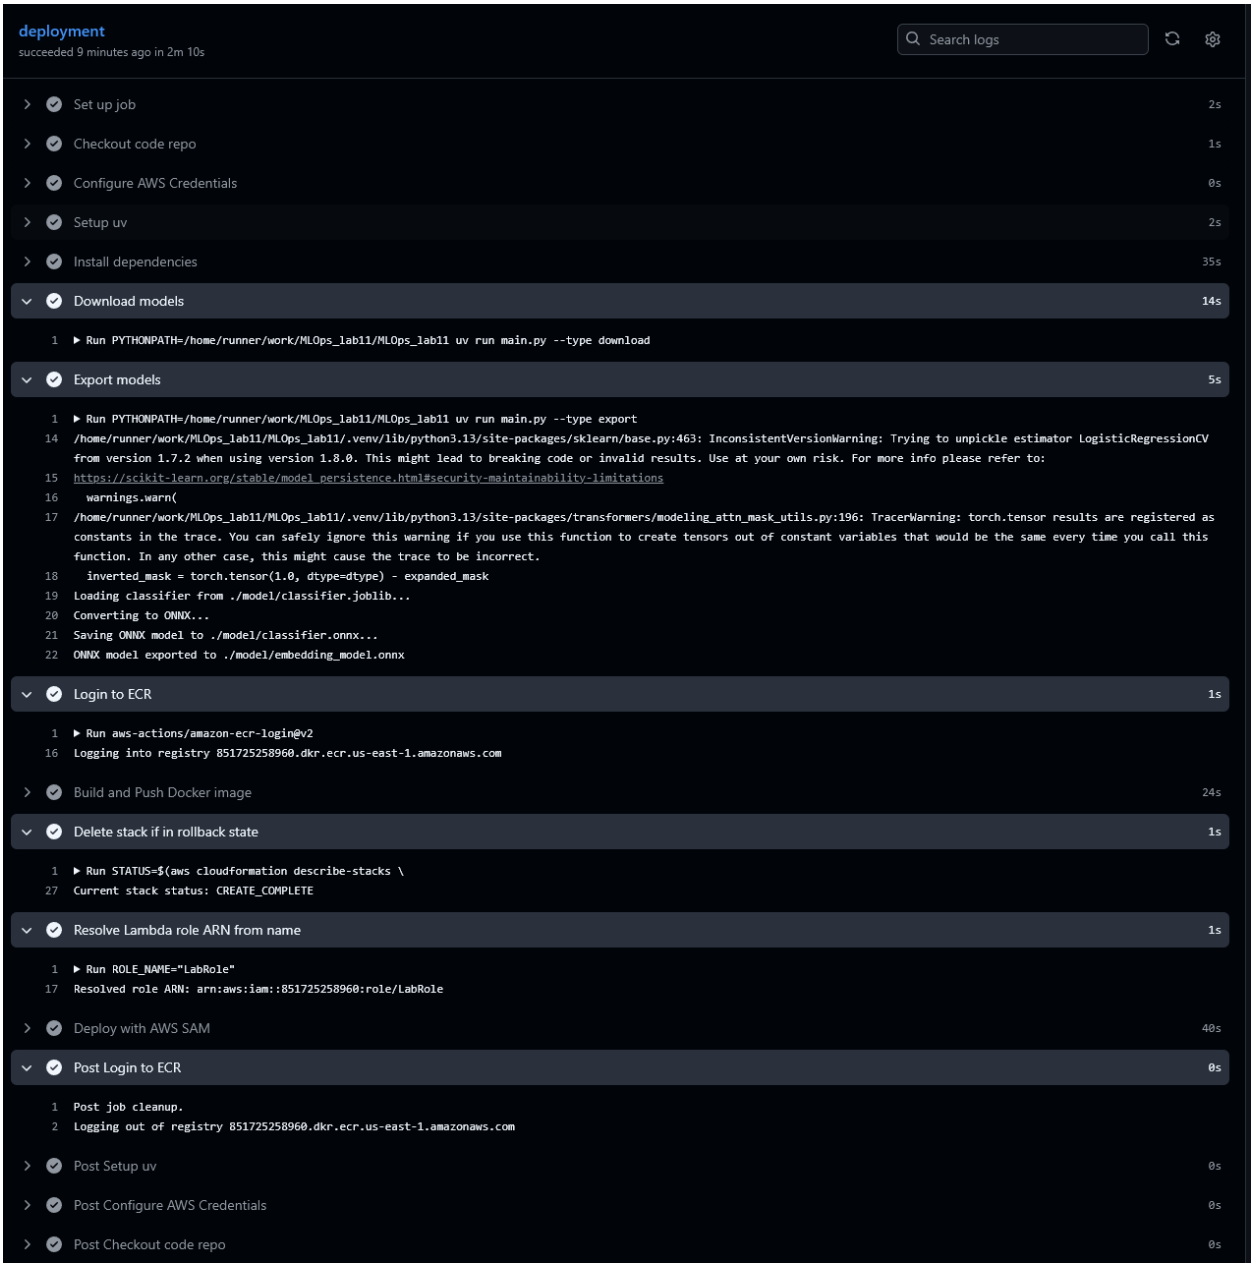
deployment
succeeded 9 minutes ago in 2m 10s
Search logs
Set up job 2s
Checkout code repo 1s
Configure AWS Credentials 0s
Setup uv 2s
Install dependencies 35s
Download models 14s
1 ▶ Run PYTHONPATH=/home/runner/work/MLOps_lab11/MLOps_lab11 uv run main.py --type download
Export models 5s
1 ▶ Run PYTHONPATH=/home/runner/work/MLOps_lab11/MLOps_lab11 uv run main.py --type export
14 /home/runner/work/MLOps_lab11/MLOps_lab11/.venv/lib/python3.13/site-packages/sklearn/base.py:463: InconsistentVersionWarning: Trying to unpickle estimator LogisticRegressionCV from version 1.7.2 when using version 1.8.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to:
15 https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
16 warnings.warn(
17 /home/runner/work/MLOps_lab11/MLOps_lab11/.venv/lib/python3.13/site-packages/transformers/modeling_attn_mask_utils.py:196: TracerWarning: torch.tensor results are registered as constants in the trace. You can safely ignore this warning if you use this function to create tensors out of constant variables that would be the same every time you call this function. In any other case, this might cause the trace to be incorrect.
18 inverted_mask = torch.tensor(1.0, dtype=dtype) - expanded_mask
19 Loading classifier from ./model/classifier.joblib...
20 Converting to ONNX...
21 Saving ONNX model to ./model/classifier.onnx...
22 ONNX model exported to ./model/embedding_model.onnx
Login to ECR 1s
1 ▶ Run aws-actions/amazon-ecr-login@v2
16 Logging into registry 851725258960.dkr.ecr.us-east-1.amazonaws.com
Build and Push Docker image 24s
Delete stack if in rollback state 1s
1 ▶ Run STATUS=$(aws cloudformation describe-stacks \
27 Current stack status: CREATE_COMPLETE
Resolve Lambda role ARN from name 1s
1 ▶ Run ROLE_NAME="LabRole"
17 Resolved role ARN: arn:aws:iam::851725258960:role/LabRole
Deploy with AWS SAM 40s
Post Login to ECR 0s
1 Post job cleanup.
2 Logging out of registry 851725258960.dkr.ecr.us-east-1.amazonaws.com
Post Setup uv 0s
Post Configure AWS Credentials 0s
Post Checkout code repo 0s

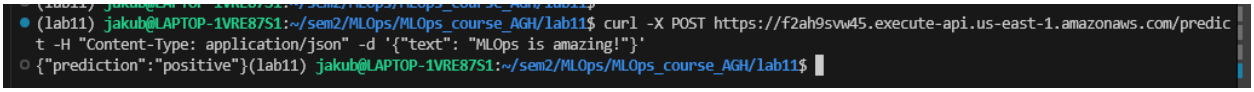
(lab11) jakub@LAPTOP-1VRE87S1:~/sem2/MLOps/MLOps_course_AGH/lab11$ curl -X POST https://f2ah9svw45.execute-api.us-east-1.amazonaws.com/predict -H "Content-Type: application/json" -d '{"text": "MLOps is amazing!"}'
{"prediction":"positive"}(lab11) jakub@LAPTOP-1VRE87S1:~/sem2/MLOps/MLOps_course_AGH/lab11$